

Final Report

Help Boost Our Online Reach

# ML PROJECT

TEAM - WeSmurfing

IMT2019073 Ronit Jain

IMT2019516 Rakshit Gupta

# Problem Statement

The task is to identify the relevant, high-quality web pages from a pool of user-curated web pages to identify “**ad-worthy**” web pages. The challenge required us to build large-scale, end-to-end **machine learning models** that can classify a website as either “**relevant**” or “**irrelevant**”, based on attributes such as **alchemy category and its score**, meta-information of the web pages, and a **one-line description** of the content of each webpage. To facilitate this, the agency has created a dataset of raw Html, meta statistics, and a binary label for each webpage. The binary label represents whether the webpage was selected for ad placement or not.

# About the Dataset

The data set was provided on Kaggle. The data set included three files:

- train.csv - the **training set** containing the target variable '**label**.'
- test.csv - the **test set**, which does not contain the target variable 'label.'
- html\_content.zip - A zip file containing the **raw HTML content** for each URL

The dataset contains an **ID column**, **25 features** and **one target *label* column**.

This dataset includes rows of web pages and their description and meta-statistics, along with a label classifying them as either "**ad-worthy**" or **not "ad-worthy"**, in **binary values 0 and 1**, respectively.

The **training data** set contained **5916** data points.

The **test data** set contains **1479** data points.

# Exploratory Data Analysis

## Countplot

---

Displayed the number of Label counts.

## Outliers

---

Displayed boxplot to find outliers in the data.

## Correlation Map

---

Created heatmap for the rest numerical features.

# Countplot Label Counts



```
print(df_train["label"].value_counts())  
sns.countplot(x='label',data=df_train)
```

[246] ✓ 0.1s

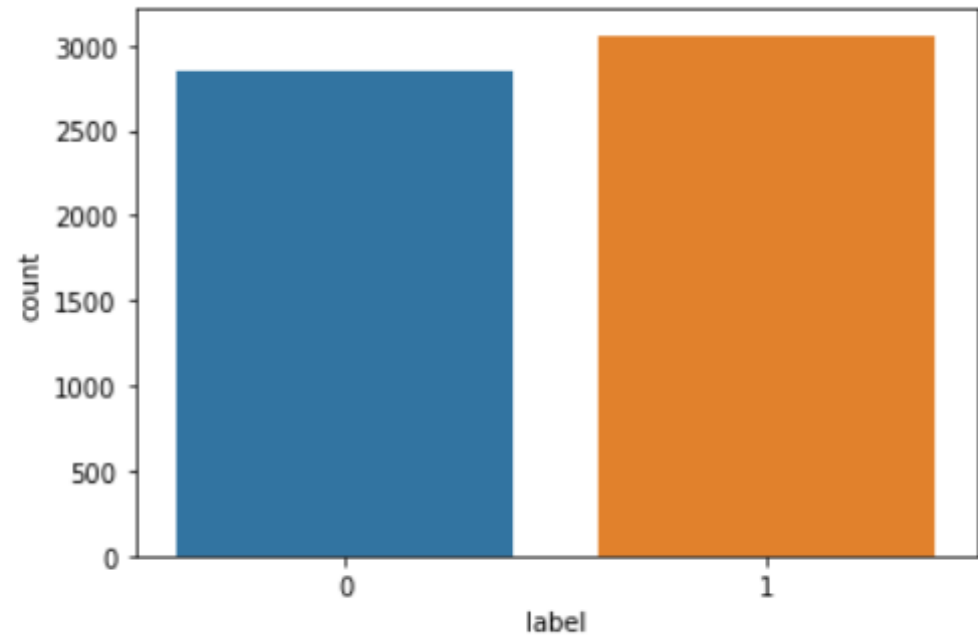
... 1 3060

0 2856

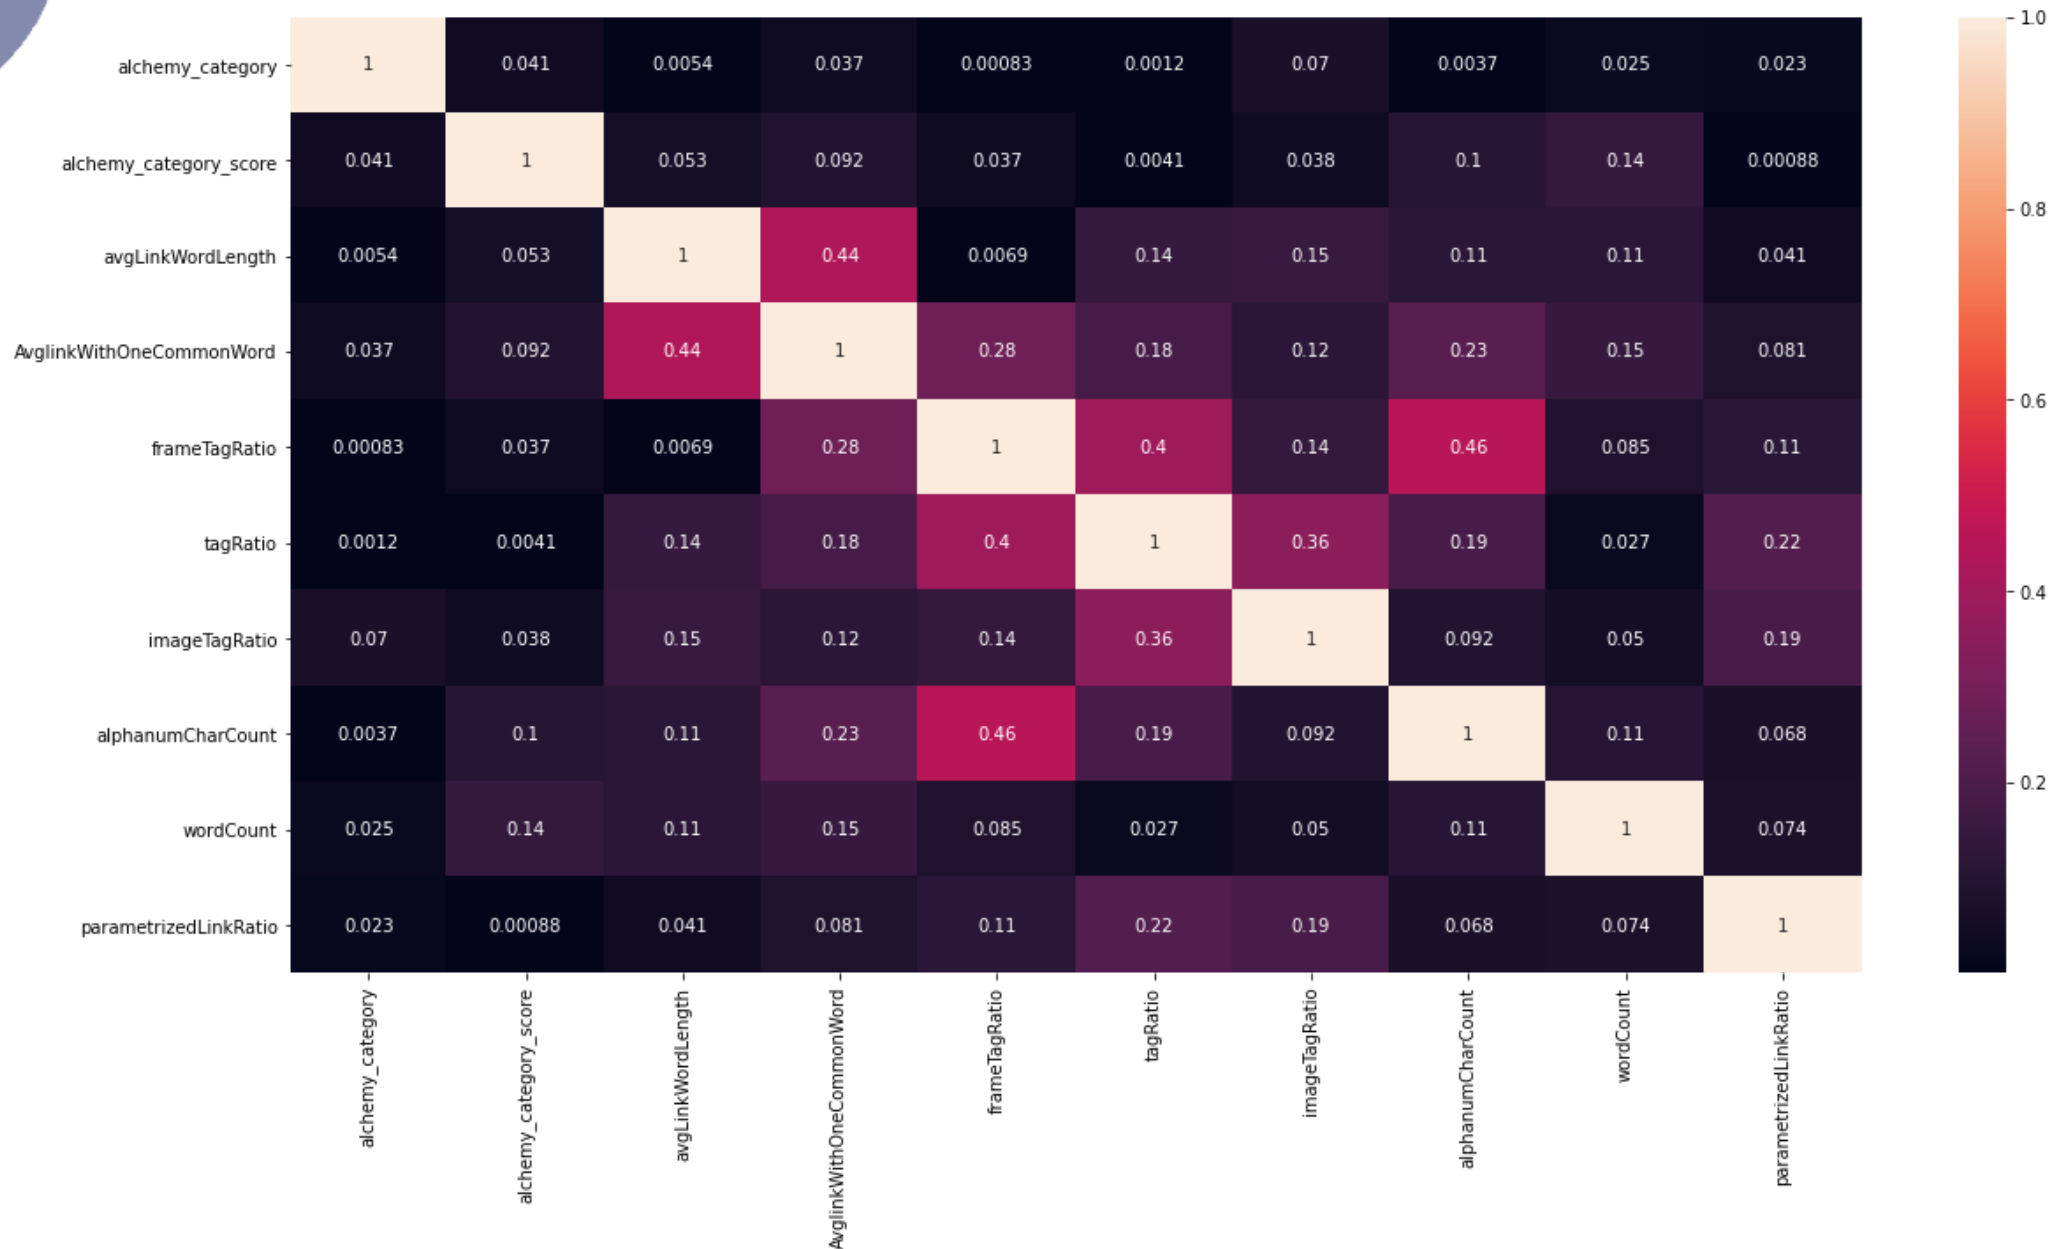
Name: label, dtype: int64

<AxesSubplot:xlabel='label', ylabel='count'>

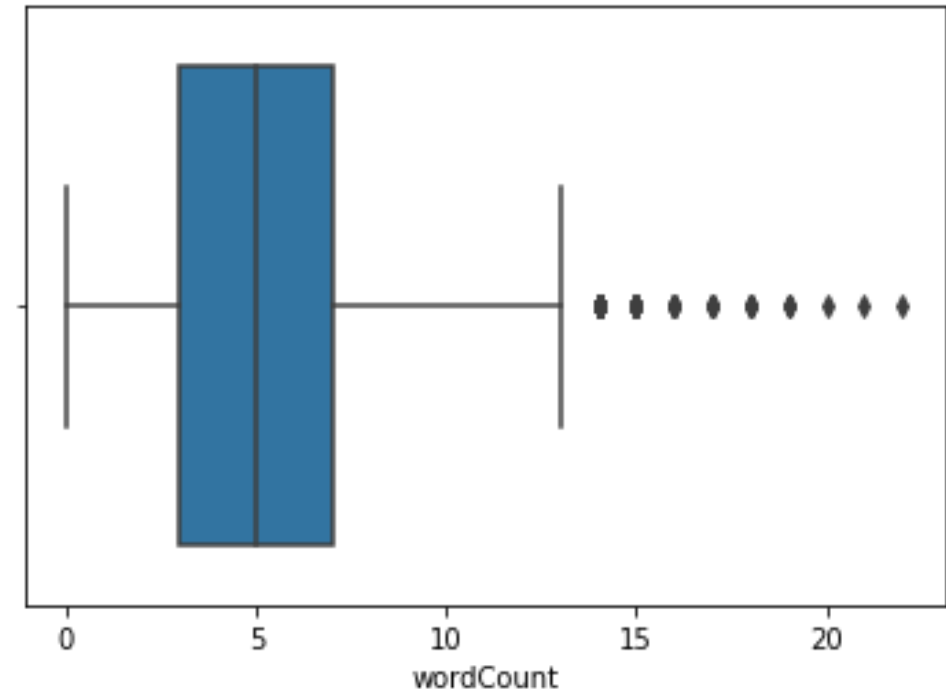
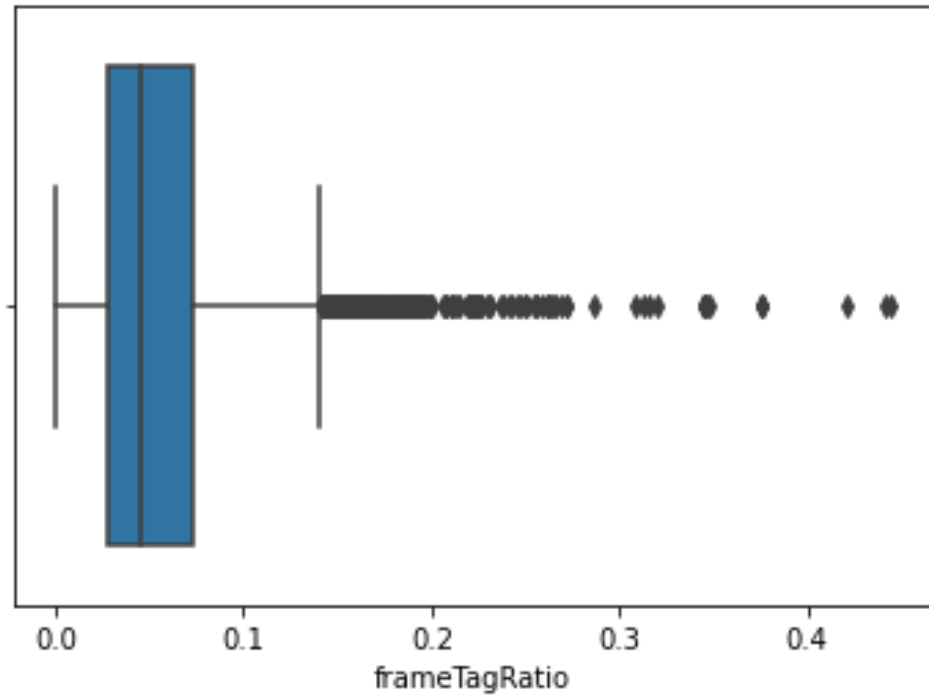
</>



# Correlation Heatmap



# Outliers



... also for many more features

# Preprocessing

- The feature "**webpageDescription**" has lots of data in text format, which will eventually be converted into numerical data.
- The feature named "**url**" is **dropped** because it's of no use in predictions.
- Removed the features "**framebased**" and "**isNews**" because they have only one kind of data.
- The **missing data** is presented in the next slide for both **train** and **test data** combined.



# Missing Data

	Missing_Data	Percentage
isNews	2843	38.444895
alchemy_category	2342	31.670047
alchemy_category_score	2342	31.670047
isFrontPageNews	1248	16.876268

- All **NaN** values are replaced with "**unknown**" in "**alchemy\_category**", **median** in "**alchemy\_category\_score**", and **mode** in "**isFrontPageNews**".

# Outliers Removal

- We used the following method to deal with the outliers.
- In the following function, if there's an outlier on either side of the whiskers, then that point is reset at the nearest whisker point.

```
class OutlierRemoval:
    def __init__(self, lower_quartile, upper_quartile):
        self.lower_whisker = lower_quartile - 1.5*(upper_quartile - lower_quartile)
        self.upper_whisker = upper_quartile + 1.5*(upper_quartile - lower_quartile)
    def removeOutlier(self, x):
        return (x if x <= self.upper_whisker and x >= self.lower_whisker else (self.lower_whisker if x < self.lower_whisker else (self.upper_whisker)))
```

```
for i in range(1, 21):
    column = df.iloc[:,i]
    # print(column.name)
    column_outlier_class = OutlierRemoval(column.quantile(0.25), column.quantile(0.75))
    outlier_removed_column = column.apply(column_outlier_class.removeOutlier)
    df[column.name] = outlier_removed_column
```

# Feature Selection and Engineering

- **After** dealing with **outliers**, there were some features with the **same** kind of **data**, like "**embedRatio**", "**domainLink**", and "**isFrontPageNews**", so we **dropped** them.
- We used **correlation** and plotted **heatmap** to fetch out the **valuable features**. Also, to remove the **highly correlated** features.
- We also used **PCA** for **feature reduction**, but it didn't give better results, so we **dropped** it.
- We used the **vectorization** method to **create features** for the text data in "**webpageDescription**".

# NLP

```
def nlp_func(text):  
    text = "".join([char for char in text if char not in string.punctuation])  
    tokens = re.split('\W+', text)  
    # text = [ps.stem(word) for word in tokens if word not in stopwords]  
    text = [wn.lemmatize(word) for word in tokens if word not in stopwords]  
    return text
```

- We used the above methods like **removing punctuation**, **tokenization**, **removing stopwords** and **lemmatization** on the text data of the feature "**webpageDescription**".
- We also tried it using **stemming** instead of **lemmatization** but didn't get better results.

# Vectorization

- We used the **TF-IDF Vectorizer** to convert the text into numeric data and **create features**.
- This process rules out more than **1lakh features**, so we used some **hyperparameters** as shown below to **limit** the no. of **features** to **< 50k**.
- We used the above-discussed **nlp\_func** for the **analyzer** part.

```
tfidf_vect = TfidfVectorizer(sublinear_tf=True, strip_accents='unicode', tokenizer=lambda x: re.findall(r'^\p{p}\W+', x), analyzer=nlp_func,  
                             token_pattern='(?u)\b\w+\b\w{,1}', min_df=5, norm='l2', ngram_range=(1,2))  
  
x_tfidf = tfidf_vect.fit_transform(df['webpageDescription']).toarray()  
x_tfidf.shape
```

- We used "**StandardScaler**" to scale the rest **10 features** of the dataset and then **merged** them with the **features** received **after vectorization**.

# Models Used

## 1. Logistic Regression Model

```
LR_model = LogisticRegression(penalty='l2', dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None,  
                               solver='lbfgs', max_iter=10000, multi_class='auto', verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)  
LR_model.fit(train_X, train_y)
```

## 2. Random Forest Model

```
RF_model = RandomForestClassifier(n_estimators=1000, criterion='gini', max_depth=None, min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0,  
                                 max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0, bootstrap=True, oob_score=False, n_jobs=None,  
                                 random_state=None, verbose=0, warm_start=False, class_weight=None, ccp_alpha=0.0, max_samples=None)  
RF_model.fit(train_X, train_y)
```

## 3. Naive Bayes Model

```
NB_model = GaussianNB()  
NB_model.fit(train_X, train_y)
```

# Scores

(Area under the Curve)

Logistic Regression Model	0.8820
Random Forest Model	0.8691
Naive Bayes Model	0.7328

# Conclusion

- This project helped us gently get the idea of the domain of NLP.
- We used it to convert the string attributes of the dataset to some form of numerical data and then constructed ML models on this numerical data.
- We also faced some challenges; eventually, those problems only helped us learn new things and apply them.



# Individual Contributions and References

## Individual Contributions

We did the whole project as a team. We discussed everything and implemented everything together.

## References

<https://towardsdatascience.com/natural-language-processing-nlp-for-machine-learning-d44498845d5b>

<https://towardsdatascience.com/text-classification-in-python-dd95d264c802>